

BANNUSHA SHAIK

Data Engineer | Python · SQL · ETL · Data Modeling · PostgreSQL
Bangalore, India · +91 8074671779 · bannushashaik85@gmail.com · linkedin.com/in/bannushashaik400 · bannusha.com

SUMMARY

Data Engineer for operational data models, pipelines, and the tables a team can query. Fortune 500 intern at SKF: Microsoft 365, SharePoint, Power Apps, and Power BI score-cards that flag a result against a qualifying bar. Built snapshot pipelines and schema-locked SQL so a number is defended against a baseline, not a raw total. Seeking a Data Engineer role.

PROFESSIONAL EXPERIENCE

Data Engineer Intern, SKF (Fortune 500 Industrial)

May 2026 – Present | Bangalore

Microsoft 365 · SharePoint · Power Apps · Power BI · operational data models

- Own the operational data routine on the SKF tenant: replaced paper 5S plant audits with an 11-screen Power Apps and SharePoint system factory teams use for scores, photos, and actions.
- Modeled the grain as one audit across five pillars, with identity-aware, zone-scoped write-back, so every submitted audit is a row plant leads can review.
- Served score-card and compliance-trend views: per-pillar totals, pass/fail against a qualifying bar, and the weakest pillar flagged for leadership.

GenAI Data Analytics Job Simulation, Tata Group (Forage)

2024

- Modeled delinquency risk on structured financial datasets and wrote the next action a stakeholder could take. A job simulation.

Technical Support, IEEE Student Branch, PES College of Engineering

Aug 2023 – Sep 2024

- College club role. Supported technical sessions: setup, troubleshooting, and keeping events running.

PROJECTS

Creator Lab | snapshot pipeline and scoring table

know-stats.streamlit.app

Python · SQLite · YouTube Data API · pandas · Streamlit · 100 creators · about 3,500 videos

- Ingested the YouTube Data API into SQLite. Grain is one snapshot of one video, so views, likes, and comments are a time series rather than a one-time scrape.
- Defined engagement as (likes + comments) / views and outlier as views versus that creator's own median, so a small channel is not ranked against a large one.
- Shipped a public app that scores a draft title and cut against the table and returns a view range plus the nearest videos already stored.

Schema-Locked SQL | English to SQL across 8 dialects

GitHub

SQL · PostgreSQL · Node.js · parameterized queries · JWT

- Built an English-to-SQL service for PostgreSQL, MySQL, SQLite, SQL Server, BigQuery, Snowflake, Oracle, and DuckDB. A pasted CREATE TABLE is the allowlist, so generation cannot invent a column. Query history is stored in PostgreSQL.

Bengaluru Urban Warehouse | city feeds into one model

In progress

Python · PostgreSQL · PostGIS · weather, air, and mobility

- Designing a warehouse for three public city feeds that do not share a schema, a clock, or a location. Each pull lands untouched. Invalid rows go to a reject table. Ward, station, and date are dimensions. An unknown location is surrogate key -1.

TECHNICAL SKILLS

Data engineering: ETL, data pipelines, data modeling, API ingestion, schema design, snapshot tables, data quality

SQL & databases: SQL, PostgreSQL, MySQL, SQLite, joins, window functions, parameterized queries

Python: pandas, NumPy, API clients, scikit-learn

Tools: Power BI, Streamlit, Grafana, Docker, Git, Power Apps, SharePoint, Microsoft 365

EDUCATION & CERTIFICATIONS

B.Tech, Artificial Intelligence & Machine Learning, PES College of Engineering

2023 – Expected 2027

Machine Learning Specialization, Coursera · IBM Generative AI Prompt Engineering · Tata Group GenAI Data Analytics (Forage), 2024