

BANNUSHA SHAIK

Data Engineer | Python · SQL · ETL · Data Modeling · PostgreSQL

Bangalore, India · +91 8074671779 · bannushashaik85@gmail.com · linkedin.com/in/bannushashaik400 · bannusha.com

SUMMARY

Data Engineer for operational models, pipelines, and the tables a team can query. Fortune 500 intern at SKF, Bangalore: paper 5S audits are now an 11-screen Power Apps and SharePoint system factory teams use, with Power BI score-cards that flag a result against a qualifying bar. Independently shipped a snapshot pipeline of 100 creators and about 3,500 videos, and an English-to-SQL service locked to the pasted schema. A city warehouse is in progress.

PROFESSIONAL EXPERIENCE

Data Engineer Intern, SKF (Fortune 500 Industrial)

May 2026 – Present | Bangalore

SQL · data modeling · ETL · data pipelines · schema design · data quality · Power BI

- Replaced paper 5S plant audits with an 11-screen Power Apps and SharePoint system. Factory teams use it for scores, photos, and actions, not files in email.
- Set the grain as one audit across five pillars. Write-back is identity-aware and zone-scoped, so a submitted audit is a row plant leads can review.
- Built the score-card and Power BI views: per-pillar totals, pass/fail against the qualifying bar, and the weakest pillar flagged for leadership.

GenAI Data Analytics Job Simulation, Tata Group (Forage)

2024

- Worked delinquency risk on structured financial datasets and wrote the next action a stakeholder could take, tied to the fields in the data. A job simulation, not employment.

Technical Support, IEEE Student Branch, PES College of Engineering

Aug 2023 – Sep 2024

- College club. Setup and troubleshooting for technical sessions so the room was ready before the talk started.

PROJECTS

Creator Lab | snapshot pipeline and scoring table

Python · SQLite · YouTube Data API · pandas · scikit-learn · Streamlit · 100 creators · about 3,500 videos

- Pulled the YouTube Data API into SQLite. Grain is one snapshot of one video, across 100 Indian creators (50 tech, 50 fashion) and about 3,500 videos.
- Engagement is (likes + comments) / views. Outlier is views divided by that creator's own median, so a small channel is not ranked against a large one.
- Streamlit scores a draft and returns a view range plus the nearest videos already stored. pandas and scikit-learn sit on the feature table.

Schema-Locked SQL | English to SQL across 8 dialects

Node.js · PostgreSQL · parameterized execution · JWT · 8 SQL dialects

- English to SQL for PostgreSQL, MySQL, SQLite, SQL Server, BigQuery, Snowflake, Oracle, and DuckDB. A pasted CREATE TABLE is the allowlist, so generation cannot invent a column.
- Query history is stored in PostgreSQL. Execution is parameterized, and the API sits behind JWT, so a generated statement cannot invent a column or concatenate input into SQL.

Bengaluru Urban Warehouse | city feeds into one model

In progress

Python · PostgreSQL · PostGIS · weather, air, and mobility

- Not shipped. Three public feeds — weather, air, and mobility — that do not share a schema, a clock, or a location. Each raw payload is kept.
- Invalid rows go to a reject table with the rule that failed. Ward, station, and date are dimensions in PostgreSQL and PostGIS. An unknown location is surrogate key -1, not a dropped row.

TECHNICAL SKILLS

Data engineering: ETL, data pipelines, data modeling, API ingestion, schema design, snapshot tables, data quality

SQL & databases: SQL, PostgreSQL, MySQL, SQLite, joins, window functions, parameterized queries

Python: pandas, NumPy, API clients, scikit-learn

Serve & plant: Power BI, Streamlit, Grafana, Docker, Git, Power Apps, SharePoint, Microsoft 365

EDUCATION & CERTIFICATIONS

B.Tech, Artificial Intelligence & Machine Learning, PES College of Engineering

2023 – Expected 2027

Machine Learning Specialization, Coursera · IBM Generative AI Prompt Engineering · Tata Group GenAI Data Analytics (Forage), 2024